

XORZEN AI Training Data Valuation Report

14-Phase Comprehensive Analysis
Licensing Value for AI Model Training

Repository: akikfaraji/DevNet (XORZEN)

Analysis Date: September 12, 2026

Prepared by: Z.ai Autonomous Valuation Agent

Table of Contents

1. Phase 1: Complete Codebase Inventory
2. Phase 2: Provenance / Rights Audit
3. Phase 3: AI-Training Value Analysis
4. Phase 4: Training-Corpus Value Model
5. Phase 5: Market Comparables
6. Phase 6: Valuation
7. Phase 7: License Structure and Price
8. Phase 8: Exclusivity Premium
9. Phase 9: \$1,000 Test
10. Phase 10: Buyer Research
11. Phase 11: Zero-Budget Path
12. Phase 12: Outreach Package
13. Phase 13: Negotiation Strategy
14. Phase 14: Final Valuation Report

Phase 1: Complete Codebase Inventory

1.1 Repository Overview

Metric	Value
Repository	akikfaraji/DevNet (XORZEN)
Project	XorZen - Hybrid Transformer x MoE x SSM Framework
Version	0.2.4 (architecture v0.4)
Author	Akik Faraji / Fraziym Tech
License	XORZENX Proprietary License (NOT open-source)
Python Requirement	>= 3.10
Development Period	2026-08-15 to 2026-09-02 (18 days)
Total Commits	36
Contributors	4 identities (likely 1 person)

1.2 File Count and Breakdown

File Type	Count	Notes
.py	162	Source, tests, scripts, benchmarks
.json	45	Reports, tokenizer vocabularies, audit results
.md	26	Documentation
.csv	84	Training metrics (all in logs/)
.log	52	Training run logs
.pdf	2	Forensic audit + benchmark reports
.docx	1	Architecture analysis document
.jsonl	1	Training corpus (data/corpus.jsonl)
Total (excl. .git)	~432	

1.3 Lines of Code by Language

Language/Category	Lines
Python - xorzen/ source (excl. pretrained JSON)	40,207
Python - tests/	3,508
Python - scripts/	23,014
Python - benchmarks/	1,650
Python - total	68,379
JSON (pretrained tokenizers)	374,232
JSON (reports/audit)	17,625
Markdown (all .md)	7,545
Grand total (all files)	~496,583

1.4 Core Library Breakdown (xorzen/)

The core library contains the most valuable source code from an AI-training perspective. The flagship component is the zeroModel architecture, which orchestrates the hybrid Transformer-MoE-SSM design. Key novel contributions include the 4-axis AdaptiveRouter, SlicedFFN with genuine per-token width sparsity, disk-sharded MoE with LRU cache, and the 6-component latent Chain-of-Thought vector system. Standard components include SSM kernels adapted from Mamba/S4, Switch Transformer load balancing, and BPE tokenization via HuggingFace libraries.

Module	Approx Lines	Classification
config.py (ConfigFactory, ModelConfig)	2,278	B - Likely original
model/components/hass_block.py	1,441	B/D - Original arch, adapted SSM
model/components/routing.py	1,698	A - Clearly original (4-axis)
model/zmoe.py (ShardedExpertFabric)	894	A - Clearly original
model/components/sliced_ffn.py	255	A - Clearly original (novel)
model/components/cot_vector.py	467	A - Clearly original (6-comp CoT)
model/ssm.py (S4DKernel)	288	D - Derivative of Mamba/S4
models/zero/model.py	1,614	B - Original orchestration
utils/sppq.py (quantization)	1,997	B - Custom quantization
utils/math_utils.py	1,304	B - Mixed standard + custom
tokenizer/ (all)	~4,100	C - Wraps HuggingFace tokenizers
data/ (all)	~6,500	B - Custom pipeline
training/ (all)	~3,600	B - Standard training loop

1.5 Documentation Volume

The repository contains approximately 7,545 lines of Markdown documentation across 26 files. The primary documents include the README (566 lines), FRAZIYM_PROJECT_SPECS_MASTER (1,745 lines), rebuild roadmaps, forensic audit reports, adversarial audit reports, and pretraining readiness assessments. While substantial in volume, much of this documentation is self-referential audit material rather than externally useful technical documentation. The architecture documentation describes the hybrid design but does not include formal specifications, API references in the traditional sense, or independent benchmarks against external baselines beyond what is in the README.

1.6 Test Coverage

The test suite comprises approximately 3,508 lines across 12 test files, containing roughly 59 distinct tests. Tests cover architecture correctness, sparsity dispatch, regression, SSM/MoE/routing correctness, blocker fixes, CoT consistency, causal leakage, load balance, SPPQ schedule, and tokenizer roundtrip. This is a reasonable test suite for the codebase size, though 59 tests for 40K+ lines of core code represents relatively thin coverage. The test suite contains the notable 137/138 pass rate issue where the remaining causal leakage test in training mode was previously dismissed as a 'known limitation.'

1.7 Component Classification Summary

Classification categories: A = Clearly original, B = Likely original (needs verification), C = Third-party, D = Derivative/adapted, E = Unknown provenance, F = Potentially unsuitable for commercial licensing.

Category	Components	% of Core
A - Clearly original	AdaptiveRouter, SlicedFFN, ShardedExpertFabric, CoT vector, disk sharding	~25%
B - Likely original	ConfigFactory, zeroModel, SPPQ, math_utils, data pipeline, training	~45%
C - Third-party	Tokenizer (HuggingFace wrapper), train_superintelligence (GPT-2)	~15%
D - Derivative	SSM/S4D kernel (from Mamba), load balance (from Switch Transformer)	~10%
E - Unknown	Some utility code, AI-assisted docstrings	~5%
F - Unsuitable	Pretrained tokenizer JSONs, runtime logs, AI-generated boilerplate	N/A (non-core)

1.8 Critical Finding: Already Public on GitHub

The entire XORZEN repository is publicly accessible on GitHub at akikfaraji/DevNet. This is the single most important fact for valuation. AI companies can already access, crawl, and use this code under GitHub's Terms of Service for model training. Public repositories have been used to train models like GitHub Copilot, Codex, and CodeLlama without payment to repository owners. The licensing value of publicly available code is effectively zero unless the owner can offer something beyond mere access, such as explicit legal clearance, curated/annotated versions, or exclusive rights to derivative training data.

Phase 2: Provenance / Rights Audit

2.1 License Analysis

CRITICAL: The XORZENX Proprietary License explicitly prohibits the very activity this valuation examines. Section 2(i) of the license states that users may not 'Extract or replicate the HASS, MoE, or routing mechanisms.' This directly targets AI training use. The license also prohibits redistribution, derivative works, commercial use, and sublicensing without written permission.

This means that **even if a buyer wanted to license this code for AI training, the current license structure blocks it.** The owner would need to issue a separate AI-training license or modify the existing license to permit this use. Until that happens, the codebase is legally unusable for AI training by any third party.

2.2 License Contradiction

The formal LICENSE file is proprietary, but the PKG-INFO (METADATA) contains contradictory statements. It claims 'It's MIT License. I think that means you can do whatever you want with it' and admits the code was made with 'some AI help (it's a secret!)'. This creates legal ambiguity. A potential buyer would need the owner to clarify and resolve this contradiction before any license could be granted. The informal MIT claim in PKG-INFO does not override the formal proprietary LICENSE file, but it could create confusion about the owner's actual intent.

2.3 Patent Claims

The license asserts: 'Certain features and innovations in XORZENX are subject to pending patent applications.' No patent numbers, application numbers, or jurisdictions are specified. A potential buyer would need to know: (a) whether patents actually exist or are merely aspirational claims, (b) what jurisdictions they cover, (c) what specific claims they contain. Patent claims could increase value if they are real and enforceable, but they could also deter buyers if they create infringement risk. Without verifiable patent information, this claim is a liability, not an asset.

2.4 AI-Generated Content

The owner has admitted that portions of the codebase were created with AI assistance. Evidence includes: (1) the PKG-INFO statement 'Made this with lots of coffee, late nights, and some AI help (it's a secret!)', (2) uniformly structured, verbose docstrings across all modules consistent with AI generation patterns, (3) 'Production-grade' appearing in virtually every module docstring, and (4) explicit attribution in scripts to 'Generated by Super Z (Z.ai)'. AI-generated code creates two problems for AI training value: (a) potential circularity/contamination if the code was itself generated by a model the buyer is training, and (b) questions about originality and copyright ownership of AI-generated portions, which varies by jurisdiction.

2.5 Dependency Rights

All runtime and development dependencies use permissive licenses (MIT, Apache-2.0, BSD-3). No copyleft (GPL) dependencies were found. This is favorable -- it means the dependencies do not impose licensing restrictions that would prevent commercial AI-training licensing of the original code. However, the derivative components (SSM from Mamba/S4 under Apache-2.0, Switch

Transformer load balance) are already well-represented in open-source training data, reducing their incremental training value to near zero.

2.6 Contributor Ownership

The repository has 4 git identities across 36 commits, but these likely represent a single person (akikfaraji, Akik Faraji, dev, Z User). There are no CLAs (Contributor License Agreements) or contributor agreements. If any contributor was a distinct person who did not assign rights, they could claim joint ownership of their contributions. This is a standard risk for small projects but is unlikely to be a practical problem given the probable single-author nature of the repository.

2.7 Personal Information / Secrets

Personal email akikfaraji@gmail.com appears in LICENSE, METADATA, and pyproject.toml. No API keys, credentials, or secrets were found in the inspected files (previous sessions sanitized git_push.py and push.py to remove a PAT token). This is a minor issue -- the email would need to be handled appropriately in any licensed dataset, but does not block licensing.

2.8 Provenance Risk Summary

Issue	Risk Level	Blocks Licensing?	Remediable?	Cost to Fix
Proprietary license prohibits AI training extraction	CRITICAL	Yes	Yes (owner can re-license)	\$0 (owner action)
License contradiction (formal vs PKG-INFO)	HIGH	Yes (creates ambiguity)	Yes (clarify intent)	\$0 (owner action)
Unverified patent claims	HIGH	No (but deters buyers)	Partially (verify patents)	Variable
AI-generated content admitted	MEDIUM	No (but reduces value)	No (already created)	N/A
Code already public on GitHub	CRITICAL	Yes (destroys scarcity)	No (already public)	N/A
Single contributor / no CLA	LOW	No	Yes (add CLA)	\$0
PII in metadata	LOW	No	Yes (remove)	\$0
Derivative components (Mamba/S4, Switch)	LOW	No	No (already derived)	N/A

Phase 3: AI-Training Value Analysis

3.1 Source Code Quality and Complexity

The core XORZEN architecture represents genuine engineering effort combining multiple paradigms (Transformer attention, State Space Models, Mixture of Experts, adaptive routing). The code quality is variable -- some components like `AdaptiveRouter` and `ShardedExpertFabric` show thoughtful design, while others have significant AI-generated boilerplate. The technical complexity is above average for a single-developer research project, but this is a double-edged sword for training value: complex code is harder for models to learn from unless it is accompanied by clear explanations, tests, and examples.

3.2 Originality

Several components are genuinely original contributions: the 4-axis `AdaptiveRouter` (combining depth/width/pathway/expert routing), `SlicedFFN` with nested weight slicing for genuine per-token width sparsity, the disk-sharded MoE with LRU cache, and the 6-component latent Chain-of-Thought vector. These represent architectural ideas not commonly found in public codebases. However, originality for AI-training value is not the same as academic novelty -- the question is whether these patterns would improve an AI coding model's capabilities. Most AI coding models need help with common patterns, not exotic architectures. The original components represent perhaps 25% of the core library by line count.

3.3 Development History and Engineering Signal

The repository has only 36 commits over 18 days of development. This is extremely thin development history. Rich commit histories with bug fixes, refactoring, design decisions, and failed approaches provide the most valuable training signal -- they show how engineers think and solve problems. XORZEN's minimal history provides almost no such signal. There are no pull requests, no issue discussions with design rationale, no failed experiments documented in commits. The audit reports and bug-fix history in the reports/ directory provide some engineering signal (problem-diagnosis-fix sequences), but this is self-authored after-the-fact documentation rather than natural development trace.

3.4 Problem-Hypothesis-Implementation-Result Chains

The repository does contain some problem-to-result chains in the audit and fix documentation. The forensic audit identified causal mask shape bugs, SSM padding issues, and `ConfigFactory` enum lookup problems. The corresponding fixes are documented. However, these chains are: (a) relatively few in number (perhaps 5-10 significant ones), (b) self-documented rather than captured in natural development workflow, (c) focused on correctness bugs rather than design decisions. High-value training data requires many such chains with rich reasoning traces. This repository does not provide enough.

3.5 Comparison with Generic GitHub Code

The majority of the XORZEN codebase consists of patterns already abundant in public GitHub repositories: PyTorch model definitions, training loops, data loaders, configuration systems, tokenizers, logging utilities, and test suites. These patterns are already well-represented in training data for models like CodeLlama, StarCoder, and DeepSeek Coder. Only the novel architectural

components (AdaptiveRouter, SlicedFFN, ShardedExpertFabric, CoT vector) represent patterns not commonly found in public code. These novel components total roughly 3,000-4,000 lines of genuinely unique code.

3.6 Domain Specialization and Scarcity

The codebase operates in the specialized domain of hybrid Transformer-SSM-MoE architecture design. This is a relatively niche area -- most public code is either pure Transformer or pure SSM (like Mamba). A combined architecture with adaptive routing is uncommon. However, 'uncommon' does not mean 'scarce in a way that creates demand.' AI labs building coding models prioritize volume of common patterns over rarity of exotic ones. A specialized architecture like XORZEN might be interesting to a lab specifically training models for ML system design, but this is a small subset of AI training activity.

3.7 Signal-to-Noise Ratio

Of the ~68K lines of Python, approximately 30-40% is AI-generated boilerplate (uniform docstrings, template-like patterns). The pretrained tokenizer JSONs (374K lines) are not useful training data -- they are vocabulary tables. The 23K lines of scripts are heavily audit/verification oriented and AI-assisted. The 52 log files and 84 CSV metric files are runtime artifacts. The genuinely useful training material is likely limited to the ~3,000-4,000 lines of novel architecture code plus perhaps 10,000-15,000 lines of non-trivial original implementation code in the core library. This gives a signal-to-noise ratio of roughly 20-25% for the Python codebase as a whole.

Phase 4: Training-Corpus Value Model

4.1 Scoring the Corpus (0-10 Scale)

Criterion	Score	Rationale
Originality	5	Some novel components (AdaptiveRouter, SlicedFFN, disk-sharded MoE, CoT vector), but 70%+ is standard patterns or derivative
Technical Sophistication	6	Above-average for a single-developer project; hybrid architecture is genuinely complex; but much complexity is from combining existing paradigms
Data Cleanliness	4	AI-generated boilerplate, license contradictions, unverified patent claims, PII in metadata, contradictory version labels
Provenance Certainty	3	Admitted AI assistance, license contradiction, unverified patent claims, no CLAs, uncertain contributor identities
Documentation Quality	4	Substantial volume (7.5K lines) but self-referential audit material; no formal API docs, no external benchmarks, no independent validation
Development-History Richness	2	Only 36 commits over 18 days; no PRs, no issue discussions, no design rationale trace; extremely thin
Engineering Signal	3	Some bug-fix chains documented but few in number; no failed experiments, no design trade-off discussions, no refactoring history
Uniqueness	5	Hybrid architecture is uncommon; novel routing/slicing/sharding patterns; but overall design borrows heavily from known architectures
Reproducibility	4	Tests exist (59) but coverage is thin (137/138 pass); no pretrained weights; training scripts present but unvalidated at scale
Benchmark Quality	3	Self-reported benchmarks only; no independent verification; no comparison on standard benchmarks beyond README claims
Quantity	3	68K Python lines total, but only ~15-19K of genuinely useful training material; tiny compared to standard training corpora (millions of examples)
Model-Training Usefulness	3	Most patterns already well-represented in public training data; novel patterns are niche and may not align with model training objectives
Commercial Demand	2	No evidence of demand for this specific type of training data; niche architecture domain; no public precedent for individual codebase licensing
Scarcity	2	Already public on GitHub; derivative components available from original sources; novel components could be replicated from descriptions
Legal Usability	1	Proprietary license PROHIBITS AI training extraction; patent claims create additional risk; license contradiction creates ambiguity

4.2 Overall Value Assessment

Weighted Average Score: 3.1 / 10

The corpus scores below the midpoint on every criterion except technical sophistication (6) and originality/uniqueness (5). The decisive factors pulling the score down are: (1) Legal Usability at 1/10 because the proprietary license explicitly prohibits AI training extraction, (2) Scarcity at 2/10 because the code is already publicly accessible on GitHub, and (3) Commercial Demand at 2/10 because there is no evidence of market demand for this specific type of training data. Even if the legal and scarcity issues were resolved, the corpus would score approximately 4.0/10, reflecting a moderately interesting but small-scale codebase with limited training value beyond the novel architectural components.

Phase 5: Market Comparables

5.1 Documented AI Training Data Licensing Deals (2024-2026)

The following deals have been verified through multiple credible sources (Reuters, WSJ, TechCrunch, NPR, publisher SEC filings). These represent the actual public record of AI training data licensing, not speculation.

Buyer	Seller	Material	Value	Structure	Date
OpenAI	News Corp	News articles, archives	>\$250M/5yr	Non-exclusive	May 2024
Meta	News Corp	News articles	~\$150M/3yr	Non-exclusive	Mar 2026
Google	Reddit	User posts, comments	~\$60M/yr	Non-exclusive	Feb 2024
OpenAI	Reddit	User posts, comments	~\$70M/yr	Non-exclusive	2024-25
Microsoft	HarperCollins	Nonfiction books	\$5K/title	Non-exclusive, 3yr	Nov 2024
Google	Spirit Airlines	Code + data (bankruptcy)	\$10M total	One-time auction	Aug 2026
Multiple	Shutterstock	Images + video	\$104M (2023)	Various	2023
Multiple	Taylor & Francis	Academic publications	\$75M	Non-exclusive	2024
Multiple	Wiley	Academic publications	\$44M	Non-exclusive	2024

5.2 Per-Unit Pricing (Where Available)

Asset Type	Per-Unit Price	Source
Book (nonfiction, backlist)	\$5,000/title (3yr non-exclusive)	Microsoft-HarperCollins
Image (stock photo)	\$1-2/image	Shutterstock deals
Book (settlement rate)	~\$3,000/book	Anthropic lawsuit settlement
Source code / codebase	NO PUBLIC PRECEDENT	N/A

CRITICAL GAP: No documented deal exists where an AI company paid specifically to license a codebase or source code repository as training data. The closest analogue is the Spirit Airlines bankruptcy auction where Google paid \$10M for the entire digital estate including code, but this was a one-time corporate liquidation, not a licensing deal, and involved enterprise-scale assets (airline operational systems, financial data, internal communications) far beyond source code alone.

5.3 Market Size

Metric	Value	Source
Global AI training dataset market (2025)	\$3.59B	Fortune Business Insights
Global AI training dataset market (2026 proj.)	\$4.44B	Fortune Business Insights
Dataset licensing CAGR to 2034	18.8%	DataIntel
Total publisher AI licensing secured	~\$2B aggregate	everything-pr.com

5.4 Relevance to XORZEN

Every documented AI training data deal involves: (a) a large content owner with millions of items, (b) content that is proprietary/not freely available, (c) content in high-demand modalities (news, social media, academic publications, images), and (d) deals valued in the millions to hundreds of millions. XORZEN fails all four conditions: it is a single small repository, the code is publicly accessible on GitHub, it is in a niche technical domain with no demonstrated demand, and it contains roughly 15-19K lines of useful training material. The HarperCollins per-book price of \$5,000/title is the closest scale analogue, but books are curated, edited, legally clean, proprietary content in a well-established market -- conditions that do not apply here.

Phase 6: Valuation

6A. Floor Value: \$0

\$0. The floor value is zero because: (1) the code is already publicly available on GitHub, where AI companies can access it under existing Terms of Service without payment, (2) the proprietary license prohibits the very AI training extraction a buyer would need, making the corpus legally unusable in its current state, and (3) there is no documented market demand for licensing individual codebases as AI training data. A rational buyer has no reason to pay for something they can already access for free, especially when the license forbids the intended use. The floor is not '\$500' or '\$200' -- it is genuinely zero given current market conditions.

6B. Fair Market Value: \$0 - \$500

If the owner resolved the licensing issues (granted explicit AI-training permission, clarified the license contradiction, verified/removed patent claims), and if the code were removed from public access, the fair market value might reach \$200-\$500 for a non-exclusive AI-training license. This is based on: (a) the HarperCollins per-book analogue of \$5,000 for a curated, legally clean, proprietary text, adjusted downward by 90-95% because XORZEN lacks curation, has uncertain provenance, contains AI-generated material, is in a niche domain with no demand signal, and provides only ~15-19K lines of useful training material. The \$200-\$500 range reflects what a buyer might pay for convenience (pre-packaged, licensed, documented) rather than for unique access to data they cannot otherwise obtain.

6C. Recommended Asking Price: \$1,500 - \$2,500

If pursuing licensing, the initial asking price should be \$1,500-\$2,500 for a non-exclusive AI-training license. This accounts for negotiation (expect 40-60% discount from asking to closing) and reflects the maximum defensible price based on: (a) the novelty of the architectural components providing some incremental value, (b) the completeness of the codebase (architecture + tests + documentation + audit trail) as a self-contained training example, and (c) the willingness-to-pay of a buyer who specifically wants licensed, documented, legally clean code for ML architecture training. This price would only be achievable after resolving all legal barriers.

6D. High-End Value: \$5,000 - \$10,000

A specialized AI lab specifically training models for ML system design or architecture synthesis might pay \$5,000-\$10,000 if: (a) the novel components (AdaptiveRouter, SlicedFFN, ShardedExpertFabric, CoT vector) are verified as genuinely novel and not replicable from published descriptions, (b) the code is removed from public access, (c) all legal issues are resolved with written assurances, (d) the lab has an immediate need for hybrid architecture training data, and (e) the license is structured as exclusive for a limited field of use. All five conditions must hold simultaneously. The probability of all five being satisfied is low.

6E. Strategic Value: \$15,000 - \$25,000

Strategic value applies only if a buyer faces a specific, expensive training-data problem that XORZEN's material directly solves. For example, if an AI lab is building a model specifically for generating novel ML architectures and has exhausted all public hybrid architecture code, XORZEN

could fill a gap. The \$15,000-\$25,000 range reflects the cost savings of not having to: (a) design and implement equivalent novel architectures from scratch (\$10-20K in engineer time), (b) write tests and documentation (\$3-5K), and (c) verify legal rights (\$2-5K). However, this scenario is speculative -- there is no evidence that any AI lab faces this specific problem or would value this specific solution.

Phase 7: License Structure and Price

All prices assume legal issues are resolved first. Prices are for a non-exclusive license unless otherwise specified.

License Type	Rights	Recommended Price	Rationale
Non-exclusive AI-training	Use in model training, no redistribution	\$1,500 - \$2,500	Standard starting point; allows multiple buyers
Limited-purpose training	Training only for ML architecture models	\$800 - \$1,500	Restricted field reduces value but increases buyer willingness
Evaluation-only	Evaluate training utility, no actual training	\$0 (with NDA)	Free evaluation to demonstrate value; NDA protects code
Commercial training	Full commercial training + fine-tuning rights	\$2,500 - \$5,000	Broader rights justify higher price
Time-limited (1 year)	Training rights for 1 year only	\$1,000 - \$1,500	Reduced from perpetual; creates renewal revenue opportunity
Field-of-use restricted	Training only for specific application domain	\$500 - \$1,000	Most restrictive; lowest price
Exclusive license	No other AI-training licenses granted	\$8,000 - \$15,000	3-5x non-exclusive premium; but limits future revenue
Perpetual license	No time limit on training rights	\$2,500 - \$4,000	One-time payment vs. annual; buyer prefers this
Single-model license	Training data for one specific model only	\$1,000 - \$2,000	Very restricted; limits buyer's flexibility

7.1 Best Balance Recommendation

The best balance is a **non-exclusive, perpetual, field-of-use-restricted AI-training license** at \$1,500-\$2,500. This structure: (a) gets you paid, (b) retains ownership and the right to license to additional buyers, (c) allows multiple revenue streams, (d) protects future value by restricting the field (preventing the buyer from using the data to build a competing product), and (e) provides the buyer with certainty (perpetual rights). The field-of-use restriction should specify 'AI model training and evaluation' to prevent the buyer from claiming rights to redistribute or commercialize the code itself.

Phase 8: Exclusivity Premium

8.1 Exclusivity Premium Calculation

If a buyer wants exclusive AI-training rights, the premium over non-exclusive should be **3-5x**. For XORZEN, this means: Non-exclusive: \$1,500-\$2,500. Exclusive: \$5,000-\$12,500. The 3-5x multiplier is based on market practice: most AI training data deals are non-exclusive (Reddit, News Corp, HarperCollins), and exclusivity typically commands a 3-5x premium in content licensing markets. The premium reflects the buyer's desire to prevent competitors from benefiting from the same data.

8.2 Should You Sell Exclusivity?

No. Selling exclusivity would be strategically unwise for this corpus. The reasons are: (1) The potential buyer pool is extremely small (perhaps 1-3 organizations would even consider this material), so exclusivity does not meaningfully increase the price a specific buyer would pay. (2) Exclusive licensing forecloses the possibility of licensing to additional buyers, but since the buyer pool is small, the incremental revenue from additional non-exclusive licenses is also small -- however, even \$500 from a second buyer is better than \$0. (3) Exclusivity reduces the owner's leverage in future negotiations. (4) The corpus is small enough that exclusivity provides minimal competitive advantage to the buyer -- they are not gaining a unique dataset that competitors cannot approximate. The recommendation is to always offer non-exclusive licenses and only grant exclusivity if the premium is truly compelling (>\$10K).

Phase 9: \$1,000 Test

Evaluating \$1,000 as an asking price for an AI-training license on the XORZEN corpus:

Dimension	Assessment	Explanation
Relative to floor value (\$0)	Too high	Floor is zero because code is public and license blocks use
Relative to fair market value (\$0-\$500)	Too high	Market value is depressed by public availability and legal barriers
Relative to recommended asking (\$1.5K-\$2.5K)	Below target	If legal issues resolved, \$1K undervalues the corpus by 33-60%
As an anchor for negotiation	Risky	Anchoring at \$1K signals low confidence; buyer will push to \$300-500
Practical achievability	Unlikely without legal fix	No buyer will pay anything while license prohibits the use

Verdict: \$1,000 is too high given current legal barriers (the code is public and the license forbids AI training use), but too low if those barriers were resolved. Do NOT anchor negotiations at \$1,000. If you fix the legal issues and make the code proprietary (remove from GitHub), anchor at \$2,000-\$2,500 for a non-exclusive license. If you cannot or will not fix the legal issues, the realistic price is \$0.

Phase 10: Buyer Research

10.1 Organizations That Acquire AI Training Data

Organization	What They Acquire	Code Relevant?	Individual Can Submit?	Payment Info	Confidence
Opendatabay	Licensed datasets (any type)	In principle, yes	Yes (seller account)	Commission: 30% (<=\$100), declining at higher tiers	HIGH
Troveo	Video, audio, gaming, robotics, business data	Not primary focus	Yes (rights holder)	\$20M+ paid out; marketplace pricing	MEDIUM
Scale AI	Labeled/annotated data	No (annotation service)	As annotator only	\$15-50/hr for annotators	LOW (wrong model)
Hugging Face	Datasets and models	Yes (datasets hub)	Yes (upload dataset)	No payment mechanism for datasets	LOW (no payment)
Stack Overflow	Technical Q&A; (licensing program)	Q&A;, not code	Corporate only	Fee undisclosed	LOW (corporate scale)
HitCents	AI Code Corpus Licensing	Yes (explicitly)	Unknown	No public pricing	MEDIUM

10.2 AI Labs That Might Value Specialized Code

The following AI labs develop coding models and could theoretically benefit from specialized ML architecture code. However, none have a documented program for acquiring individual codebase licenses:

- **DeepSeek** -- Trains DeepSeek Coder on public code; no individual licensing program known
- **Meta (CodeLlama)** -- Trains on public GitHub code; enterprise-scale data acquisition only
- **Mistral** -- Codestral trained on public code; no individual licensing program
- **Google (Gemini/Gemma)** -- Large-scale data acquisition; Spirit Airlines deal suggests openness to code licensing at enterprise scale
- **Anthropic** -- No documented code-licensing program; primarily text-focused data acquisition
- **Cursor/Windsurf/Cody** -- AI coding assistants; train on public code; no licensing programs for individual repos

Realistic assessment: No AI lab or company currently has a documented pathway for an individual developer to license a single codebase for AI training data. The closest option is Opendatabay, where you could list the XORZEN corpus as a dataset, but there is no evidence of demand for this specific type of material on that platform.

Phase 11: Zero-Budget Path

11.1 \$0 Required (Do It Yourself)

- **Resolve license contradiction:** Write a clear, unambiguous AI-training license grant (1-2 pages). Specify exactly what rights are granted, for what purpose, under what conditions. Cost: \$0 (your time).
- **Remove PII:** Remove akikfaraji@gmail.com from metadata files or confirm it is acceptable to include. Cost: \$0.
- **Remove patent claims or verify them:** If no patent application was actually filed, remove the claim from the license. If one was filed, document the application number and jurisdiction. Cost: \$0 (research) to variable (patent attorney).
- **Create a clean dataset package:** Curate the most valuable training material (novel architecture code + tests + documentation) into a separate, well-organized directory. Add a manifest describing each file's provenance classification. Cost: \$0 (your time).
- **Write a corpus description:** 1-page technical description of what the training data contains, its novel contributions, and its provenance. Cost: \$0.
- **List on Opendatabay:** Create a seller account and list the curated dataset. Cost: \$0 (commission only on sale).
- **Approach potential buyers directly:** Email AI labs and data platforms with the corpus description and asking price. Cost: \$0.

11.2 Optional Costs

- **Attorney review of license:** \$500-\$2,000 for a technology attorney to review the AI-training license grant. Recommended but not strictly necessary for initial outreach.
- **Copyright registration:** \$65-\$250 per work (US Copyright Office). Provides legal benefits in infringement cases but not required for licensing.
- **Professional data curation:** \$1,000-\$5,000 for a data engineering service to clean, annotate, and format the corpus for maximum training value. Unlikely to be worthwhile given the small corpus size.

11.3 Buyer-Triggered Costs

- **Detailed due diligence package:** Only prepare if a serious buyer requests it. May require attorney review (\$1-3K).
- **Make code private on GitHub:** If a buyer requires the code not be publicly available, you would need to make the repo private. This is free but may conflict with other goals (open-source community, academic visibility).
- **Custom license drafting:** If a buyer wants specific terms, have an attorney draft the agreement (\$2-5K). Only worthwhile if the deal value exceeds \$5K.

Phase 12: Outreach Package

A. Corpus Description (One Paragraph)

The XORZEN corpus is a self-contained Python codebase (~68K lines) implementing a hybrid Transformer-SSM-MoE architecture with several novel components: a 4-axis adaptive router (depth/width/pathway/expert), per-token width-sparse SlicedFFN, disk-sharded MoE with LRU cache, and a 6-component latent Chain-of-Thought vector. The corpus includes the core architecture, configuration system, training infrastructure, data pipeline, tokenizer, test suite (59 tests), documentation, audit reports, and development artifacts. The novel architectural components (~3-4K lines) represent patterns not commonly found in public codebases and may provide incremental value for models training on ML system design or hybrid architecture synthesis.

B. Technical Description

XORZEN v0.4 is a hybrid neural architecture combining local attention, state space models (SSM/S4D), mixture of experts (MoE), and adaptive routing. The zeroModel flagship orchestrates these components through a HASS (Hybrid Attention-SSM-Structured) block architecture. Key technical features: (1) LowRankGlobalPathway with causal masking for efficient long-range attention, (2) SSMPathway with zero-order hold discretization and parallel scan, (3) AdaptiveRouter with Gumbel-softmax routing across four axes, (4) ShardedExpertFabric with disk-based expert storage and LRU cache, (5) SlicedFFN with nested weight slicing for genuine per-token compute sparsity, (6) InternalLatentCoT with 6-component reasoning vector and auxiliary loss. Model variants range from 23K to 7B parameters. Implemented in PyTorch with no copyleft dependencies.

C. Provenance Statement

The XORZEN codebase was developed by Akik Faraji / Fraziym Tech between August 15 and September 2, 2026. The core architecture is original, with the following exceptions: SSM/S4D kernel implementation is adapted from the Mamba/S4 architecture (Gu et al., 2022; Alberti et al., 2023, Apache-2.0 licensed); Switch Transformer load balance loss follows Fedus et al. 2022; BPE tokenization uses the HuggingFace tokenizers library (Apache-2.0). All remaining code is original. Portions of the codebase (docstrings, metadata, audit scripts) were created with AI assistance. The author retains all rights not explicitly granted in the XORZENX Proprietary License.

D. AI-Training Licensing Summary

Available for non-exclusive AI model training licensing under a dedicated license grant. The license covers: use of source code, documentation, and test suite as training data for AI/ML model development and evaluation. Does not cover: redistribution of the code, commercial products derived from the code, or use of the code as software (only as training data). The novel architectural components (AdaptiveRouter, SlicedFFN, ShardedExpertFabric, CoT vector) are the primary training-value assets. Derivative components (SSM, load balance) are already available from original open-source sources.

E. Buyer Pitch

Looking for specialized ML architecture training data? The XORZEN corpus contains several novel hybrid architecture patterns -- 4-axis adaptive routing, per-token width sparsity, disk-sharded MoE,

latent CoT vectors -- that are uncommon in public codebases. Available for non-exclusive AI-training licensing at \$2,000, with full provenance documentation and legal clearance.

F. Email Template

Subject: AI Training Data Licensing Opportunity -- Novel Hybrid ML Architecture Codebase

Dear [Name],

I am writing to offer a specialized Python codebase for AI model training licensing. The XORZEN corpus implements a hybrid Transformer-SSM-MoE architecture with several novel components not commonly found in public repositories, including a 4-axis adaptive router, per-token width-sparse feedforward, disk-sharded MoE, and a 6-component latent Chain-of-Thought system. The codebase includes ~68K lines of Python, a test suite, documentation, and audit reports. I am offering a non-exclusive AI-training license at \$2,000. Full provenance documentation and a corpus manifest are available upon request under NDA. Would you be interested in evaluating this material?

Best regards, Akik Faraji

G-I. Pricing

Item	Value
Suggested Asking Price	\$2,000 (non-exclusive, perpetual, AI-training only)
Minimum Acceptable Price	\$500 (after negotiation)
Negotiation Strategy	Start at \$2,000; accept \$1,000-\$1,500; walk away below \$500

K. Due Diligence Materials

- Corpus manifest with per-file provenance classification
- Full license text and AI-training license grant
- Dependency list with all licenses verified
- AI-generated content disclosure with estimated percentages
- Test suite results (full pytest output)
- Architecture documentation (README + design docs)
- Code quality metrics (line counts, complexity, classification)
- Commit history with contributor identities
- Patent status documentation (or statement that no patents are filed)

Phase 13: Negotiation Strategy

Question	Answer
What price should I initially quote?	\$2,000 for non-exclusive AI-training license
What is my private minimum?	\$500 (below this, the transaction cost exceeds value)
When should I offer a discount?	If the buyer is a well-known AI lab (brand value of association); if they want multiple licenses; if they can provide compute/infrastructure in kind
Should I reveal my minimum price?	Never. Reveal only the asking price. Your minimum is \$500 but the buyer should not know this
Should I offer non-exclusive first?	Always. Non-exclusive preserves future licensing revenue and is easier for buyers to accept
How should exclusivity be priced?	3-5x the non-exclusive price. At \$2K non-exclusive, exclusive = \$6K-\$10K
Should I offer evaluation license first?	Yes. Free 30-day evaluation under NDA lets the buyer verify training value before committing
What evidence to show before revealing full corpus?	Corpus description, architecture overview, line counts, provenance classification, sample code snippets (not full corpus)
What should I never give away for free?	The full source code, the curated dataset, the test suite with expected outputs, or any executable version of the code

Phase 14: Final Valuation Report

Final Valuation Table

Metric	Value
Floor Value	\$0 (code is public; license prohibits AI training use)
Fair Market Value	\$0 - \$500 (if legal issues resolved)
Recommended Asking Price	\$2,000 (non-exclusive, after legal fix)
Likely Negotiated Range	\$500 - \$1,500
High-End Value	\$5,000 - \$10,000 (exclusive, specialized buyer)
Strategic Value	\$15,000 - \$25,000 (solves specific training-data gap)
Non-Exclusive License	\$1,500 - \$2,500
Exclusive License	\$5,000 - \$12,500
\$1,000 Assessment	Too high given current barriers; too low if barriers resolved

FINAL RECOMMENDATION

Primary Recommended Asking Price: \$2,000 for a non-exclusive, perpetual, AI-training-only license, contingent on first resolving the three legal barriers (license grant, license contradiction, patent claims).

Defensible Range: \$500 - \$2,500

The lower bound (\$500) reflects the minimum at which a transaction makes sense given the effort of licensing, documentation, and negotiation. The upper bound (\$2,500) reflects the maximum a buyer might pay for the convenience of a pre-packaged, legally clean, documented training corpus with novel architectural patterns. Above \$2,500, the buyer would rationally choose to either use public code or commission original code.

CONFIDENCE: LOW

Confidence is LOW because: (1) there is no public precedent for licensing an individual codebase as AI training data -- the valuation is entirely extrapolative, (2) the three legal barriers (proprietary license, license contradiction, patent claims) create a currently insurmountable obstacle to any transaction, (3) the code is already publicly accessible on GitHub, destroying the scarcity value that underpins most AI data licensing, (4) no evidence of buyer demand exists for this specific type of material, and (5) the admitted AI-generated content creates contamination/circularity concerns that reduce training value.

GO / MAYBE / NO-GO

MAYBE -- with significant prerequisites.

Licensing this material is not commercially worthwhile in its current state. The three legal barriers make the corpus unusable by any buyer. However, if the owner is willing to: (1) grant an explicit AI-training license, (2) clarify/resolve the license contradiction, and (3) either verify or remove patent claims, then licensing becomes worth pursuing -- though the expected revenue (\$500-\$1,500 after negotiation) should be weighed against the time investment of packaging, outreach, and negotiation.

The more fundamental question is whether the owner should pursue AI-training licensing at all, or instead focus on alternative monetization strategies that better leverage XORZEN's technical merits: consulting/expertise in hybrid architecture design, open-sourcing to build reputation and then offering paid support/customization, or using XORZEN as a portfolio piece to attract employment or investment. These alternatives may yield higher returns than AI-training data licensing, which is an immature market with no proven pathway for individual developers.

Summary of Every Code Change

No code changes were made as part of this valuation. This report is purely analytical. The valuation identifies what would need to change for licensing to become viable, but does not implement those changes. The required changes are legal/administrative (license modification, patent claim clarification, PII removal) rather than code modifications.

What Evidence Is Missing

- Whether patent applications were actually filed (or just aspirational claims)
- Whether the owner is willing to modify the license to permit AI training use
- Whether the owner would make the repository private (required for any scarcity value)
- Whether any AI lab or data marketplace has demand for this type of material
- Whether the AI-generated portions are contaminated by models a buyer would be training
- Verified per-unit pricing for source code as AI training data (does not exist publicly)
- Whether the novel architectural components can be replicated from published descriptions alone